Table 1: Cost analysis comparing KCEA against baselines on seven benchmarks. Each cell reports the mean ± std across datasets. Best (lowest) per column in **bold**. Alignment time is shown only for KCEA; baselines have no alignment step.

(a) Trainable parameters.

Method	**Peak train. (M)** ↓	**Final (M)** ↓	**Total (M)** ↓	**Share (%)** ↓
APER+Adapter	–	–	–	–
APER+SSF	–	–	–	–
APER+VPT-Deep	–	–	–	–
CODA-Prompt	–	–	–	–
DUCT	–	–	–	–
DualPrompt	–	–	–	–
EASE	–	–	–	–
L2P	–	–	–	–
MOS	–	–	–	–
SLCA	–	–	–	–
TUNA	–	–	–	–
KCEA	–	–	–	–

(b) Memory cost (MB).

Method	**Peak train. RAM (MB)** ↓	**Final RAM (MB)** ↓
APER+Adapter	–	–
APER+SSF	–	–
APER+VPT-Deep	–	–
CODA-Prompt	–	–
DUCT	–	–
DualPrompt	–	–
EASE	–	–
L2P	–	–
MOS	–	–
SLCA	–	–
TUNA	–	–
KCEA	–	–

All values exclude the shared frozen backbone (327 MB, ViT-B/16). For KCEA, final RAM includes class-conditional Gaussian statistics (μ_c, Σ_c).

(c) Wall-clock time (s, cumulative across all tasks).

Method	**Training (s)** ↓	**Evaluation (s)**	**Alignment (s)**
APER+Adapter	–	–	–
APER+SSF	–	–	–
APER+VPT-Deep	–	–	–
CODA-Prompt	–	–	–
DUCT	–	–	–
DualPrompt	–	–	–
EASE	–	–	–
L2P	–	–	–
MOS	–	–	–
SLCA	–	–	–
TUNA	–	–	–
KCEA	–	–	–